




Figure 8: Bias scores from each model for the two intersectional categories, separated by how the non-target differs from the target. In each case, the label that matches both dimensions (race by gender & race by SES) is the bias target and the other label is the non-target. Results are generally inconsistent across the three breakdowns for race by gender, indicating that BBQ is not measuring a significant effect of *intersectional* bias in this case.




Figure 9: By-label breakdown of race by gender bias scores from UnifiedQA. The model appears to use systematic intersectional biases in responses when that bias targets Middle-Eastern women and Asian men, but results in all other labels are more mixed.




Figure 10: By-label breakdown of race by SES bias scores from UnifiedQA. The model uses some systematic intersectional biases when the bias target is identified as being either Black/African American or Hispanic/Latinx and having low SES, but results for the other labels are more mixed.